

```
# Striving for Simplicity: The All Convolutional Net
##### *Group 16: Wei-Tse Yang, Sunwei Wang, and Qingyuan Cao*
##### 20/4/2020
In this notebook, we try to reproduce the TABLE 3 in the [original paper]
(https://arxiv.org/abs/1412.6806). The source code of one of the models on
```Pytorch``` and the training procedure on Google Colab can be found in
[Github](https://github.com/StefOe/all-conv-pytorch). **We adopt the original
training procedure and change it to the Python class. Also, we build the models
from scratch on Pytorch.**
```

---

```
Brief Introduction
The paper shows that replacing the max-pooling with the convolutional with
increased strides can improve the performance. The authors prove it by training
models with max-pooling, models removing max-pooling, and models replacing
max-pooling with convolutional layers. The results show the models replacing
max-pooling with convolutional layers with strides generally have better
performance. We provide a detailed explanation as follows. The authors tested 12
networks by designing 3 model bases and 3 branches.
```

```
Base: Model A, Model B, and Model C
Since the design of convolutional layers would influence the performance, the
authors test three model bases. **Model A** uses the 5x5 strides. **Model B**
uses the 5x5 strides but also adds one convolutional layer with 1x1 strides
after that. **Model C** uses two convolutional layers with 3x3 strides.
```

```
<img src='https://drive.google.com/uc?id=1HKDGWePX-PkBgRbeb8J_mwO-DDULhU2q'
width="600px"/>
```

```
Branch: Model, Strided-CNN, ConvPool-CNN, and ALL-CNN
Each model base has one original model and three branches. ***"Model"*** is the
model with max-pooling. ***"Strided-CNN"*** is the model removing max-pooling. **
"All-CNN"*** is the model replacing max-pooling with convolutional strides. The
better performance of "All-CNN" might result from more parameters than "Model"
and "Strided-CNN". To solve it, "ConvPool-CNN" is proposed. ***"ConvPool-CNN"***
is the model with max-pooling and one more convolutional layer before the
pooling. "ConvPool-CNN" should have the same number of parameters as "All-CNN".
Therefore, if "All-CNN" has a better performance than "ConvPool-CNN", we can
prove the better performance on "All-CNN" does not result from more parameters.
We show architecture with the base of model C in the following image.
```

```
<img src='https://drive.google.com/uc?id=1gzpTwoW_Xx8YrHZdvU0ZmFT3n-1ktx1v'
width="600px"/>
```

---

```
Experiment Setup
All 12 networks are trained on the CIFAR-10 with the stochastic gradient
descent with a fixed momentum of 0.9 and 350 epochs. The learning rate γ is
chosen from the set $\in [0.25, 0.1, 0.05, 0.01]$. It is also scheduled by
multiplying with a fixed factor of 0.1 in the epoch $S = [200, 250, 300]$. The
paper only presents the best performance among all learning rates. Based on the
source code, the performance is directly evaluated on the CIFAR-10 test set. In
```

other words, **the source code did not use the validation for hyper-parameter tuning!**

---

# Reproduction Results

Our reproduction results are different from the paper as shown in the following table. First, we perceive an error rate gap between the paper and our reproduction, which is around 5~7%. Secondly, we obtain the same ranking in **Model A**, but we fail to reproduce the same order in **Model B** and **Model C**. This is because Model B and Model C are difficult to converge than Model A. Also, **ALL-CNN-B** and **ALL-CNN-C** fail to converge to the right place. During the training, the first three learning rates seem too large for the models. Therefore, we also try another learning rate=0.001 and a longer epoch=400 in **ALL-CNN-B**, where we add the \* mark. Unfortunately, the model still fails to converge to the right place.

Model	Error Rate of Paper	Error Rate of Ours
Model A	12.47%	19.27%
Strided-CNN-A	13.46%	20.27%
<b>ConvPool-CNN-A</b>	<b>10.21%</b>	<b>15.46%</b>
ALL-CNN-A	10.30%	15.60%
<b>Model B</b>	10.20%	<b>17.01%</b>
Strided-CNN-B	10.98%	23.20%
ConvPool-CNN-B	9.33%	18.22%
<b>ALL-CNN-B</b>	<b>9.10%</b>	*29.48%
<b>Model C</b>	9.74%	<b>13.07%</b>
Strided-CNN-C	10.19%	15.49%
ConvPool-CNN-C	9.31%	14.39%
<b>ALL-CNN-C</b>	<b>9.08%</b>	17.89%

```
The example of training procedure
Choose the model A
training=Training(baseModel=[True,False,False])
Create the dataset
training.createDataset()
Choose the branch: All-CNN
training.modifiedModel=[False,False,False,True]
Start Training
training.Procedure()
```

## ▼ Valudation

Since the source code did not use the validation set for hyper-parameters tuning, we conduct the validation on the training data to create the validation set. We show the results in the following table. The performance is generally higher from the counterpart on the validation set. However, if we compare the test error with the original test error, the result reveals that the models might be overestimated.

Model	Validation Set Error	Test Error	Original Test Error
Model A	21.20%	20.45%	19.27%
Strided-CNN-A	21.72%	21.38%	20.27%
<b>ConvPool-CNN-A</b>	<b>15.93%</b>	<b>17.04%</b>	<b>15.46%</b>
ALL-CNN-A	17.57%	18.57%	15.60%
<b>Model B</b>	<b>16.65%</b>	<b>17.81%</b>	<b>17.01%</b>
Strided-CNN-B	17.53%	18.68%	23.20%
ConvPool-CNN-B	17.53%	17.51%	18.22%
ALL-CNN-B	*24.53%	*25.78%	*29.48%
<b>Model C</b>	<b>14.13%</b>	<b>14.87%</b>	<b>13.07%</b>
Strided-CNN-C	20.89%	21.67%	15.49%
ConvPool-CNN-C	17.81%	17.60%	14.39%
ALL-CNN-C	20.41%	19.13%	17.89%

```
Validation can be conducted by settting validation equal to True
training=Training(validation=True,bestModel_allLR=True,baseModel=[True,False,False])
```

## ▼ Dropout and Batch Normalization

Dropout is a simple method to prevent neural networks from overfitting. In our all convolutional net p used to regularize all networks. Dropout was almost essential in all the state-of-the-art networks before normalization(BN). With the introduction of BN, it has shown its effectiveness and practicability in the evidence that when these two techniques are used combinedly in a network, the performance of the Szegedy, 2015). In our study, we will investigate the results using BN and Dropout independently and Dropout simultaneously.

### ▼ BatchNorm *only*

```
model = Model(dropOut=False, BN=True)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 96, 30, 30]	7,296
BatchNorm2d-2	[-1, 96, 30, 30]	192
ReLU-3	[-1, 96, 30, 30]	0
MaxPool2d-4	[-1, 96, 14, 14]	0
Conv2d-5	[-1, 192, 12, 12]	460,992
BatchNorm2d-6	[-1, 192, 12, 12]	384
ReLU-7	[-1, 192, 12, 12]	0
MaxPool2d-8	[-1, 192, 5, 5]	0
Conv2d-9	[-1, 192, 5, 5]	331,968
BatchNorm2d-10	[-1, 192, 5, 5]	384
ReLU-11	[-1, 192, 5, 5]	0
Conv2d-12	[-1, 192, 5, 5]	37,056
BatchNorm2d-13	[-1, 192, 5, 5]	384
ReLU-14	[-1, 192, 5, 5]	0
Conv2d-15	[-1, 10, 5, 5]	1,930
BatchNorm2d-16	[-1, 10, 5, 5]	20
ReLU-17	[-1, 10, 5, 5]	0
AdaptiveAvgPool2d-18	[-1, 10, 1, 1]	0
Flatten-19	[-1, 10]	0
Total params: 840,606		

### ▼ BatchNorm with Dropout

```
model = Model(dropOut=True, BN=True)
```

Layer (type)	Output Shape	Param #
Dropout-1	[-1, 3, 32, 32]	0
Conv2d-2	[-1, 96, 30, 30]	7,296
BatchNorm2d-3	[-1, 96, 30, 30]	192
ReLU-4	[-1, 96, 30, 30]	0
MaxPool2d-5	[-1, 96, 14, 14]	0
Dropout-6	[-1, 96, 14, 14]	0
Conv2d-7	[-1, 192, 12, 12]	460,992
BatchNorm2d-8	[-1, 192, 12, 12]	384
ReLU-9	[-1, 192, 12, 12]	0
MaxPool2d-10	[-1, 192, 5, 5]	0
Dropout-11	[-1, 192, 5, 5]	0
Conv2d-12	[-1, 192, 5, 5]	331,968
BatchNorm2d-13	[-1, 192, 5, 5]	384
ReLU-14	[-1, 192, 5, 5]	0
Conv2d-15	[-1, 192, 5, 5]	37,056
BatchNorm2d-16	[-1, 192, 5, 5]	384
ReLU-17	[-1, 192, 5, 5]	0
Conv2d-18	[-1, 10, 5, 5]	1,930
BatchNorm2d-19	[-1, 10, 5, 5]	20
ReLU-20	[-1, 10, 5, 5]	0
AdaptiveAvgPool2d-21	[-1, 10, 1, 1]	0
Flatten-22	[-1, 10]	0
Total params: 840,606		

### ▼ Dropout *only*

```
model = Model(dropOut=True, BN=False)
```

Layer (type)	Output Shape	Param #
Dropout-1	[-1, 3, 32, 32]	0
Conv2d-2	[-1, 96, 30, 30]	7,296
ReLU-3	[-1, 96, 30, 30]	0
MaxPool2d-4	[-1, 96, 14, 14]	0
Dropout-5	[-1, 96, 14, 14]	0
Conv2d-6	[-1, 192, 12, 12]	460,992
ReLU-7	[-1, 192, 12, 12]	0
MaxPool2d-8	[-1, 192, 5, 5]	0
Dropout-9	[-1, 192, 5, 5]	0
Conv2d-10	[-1, 192, 5, 5]	331,968
ReLU-11	[-1, 192, 5, 5]	0
Conv2d-12	[-1, 192, 5, 5]	37,056
ReLU-13	[-1, 192, 5, 5]	0
Conv2d-14	[-1, 10, 5, 5]	1,930
ReLU-15	[-1, 10, 5, 5]	0
AdaptiveAvgPool2d-16	[-1, 10, 1, 1]	0
Flatten-17	[-1, 10]	0
Total params: 839,242		

### ▼ Without using BatchNorm or Dropout

```
model = Model(dropOut=False, BN=False)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 96, 30, 30]	7,296
ReLU-2	[-1, 96, 30, 30]	0
MaxPool2d-3	[-1, 96, 14, 14]	0
Conv2d-4	[-1, 192, 12, 12]	460,992
ReLU-5	[-1, 192, 12, 12]	0
MaxPool2d-6	[-1, 192, 5, 5]	0
Conv2d-7	[-1, 192, 5, 5]	331,968
ReLU-8	[-1, 192, 5, 5]	0
Conv2d-9	[-1, 192, 5, 5]	37,056
ReLU-10	[-1, 192, 5, 5]	0
Conv2d-11	[-1, 10, 5, 5]	1,930
ReLU-12	[-1, 10, 5, 5]	0
AdaptiveAvgPool2d-13	[-1, 10, 1, 1]	0
Flatten-14	[-1, 10]	0
Total params: 839,242		

We compare the results of different combination of these two techniques and generated the table below.

Model	BN only	BN + Dropout	Dropout only	No BN no Dropout
Model A	no converge	no converge	19.27%	13.82%

As shown in the table above, we used general Model A to study the two techniques Batch Normalization and Dropout. Batch Normalization increases or decreases our model performance in this case. We implemented BN layer between two ReLU activations. Dropout is also applied in between convolution layers, and it is used after the pooling layer. We used Dropout only, and we found out that using BN without Dropout or combining both BN with Dropout leads to worse performance. Xiang, et al. (2019) stated in their papers that the worse performance may be the result of variance spikes which lead the model to not converge, giving the result of 19.27%. But the performance is still not as good as 13.82% might be due to that we did not have the time to tune hyperparameter, dropout rate, we only used the dropping out inputs and 50% otherwise.

## ▼ Optimizer

The default optimizer used in the paper and in our reproduction is Stochastic Gradient Descent (SGD). We tried different optimizers since Adaptive Moment Estimation (Adam) is one of the most popular optimizers instead of SGD optimizer in our model. However, the model did not converge with Adam under the same settings even though in theory Adam combines the advantage of two SGD variants, RMSProp and AdaGrad, it is still setting to converge to an optimal solution.

SGD with momentum optimizer:

```
self.optimizer = optim.SGD(model.parameters(), lr=lr, momentum=0.9)
```

Adam optimizer:

```
self.optimizer = optim.Adam(self.model.parameters(), lr=self.lr)
```

The optimizer is the main approach nowadays for training neural networks with minimizing its error. In the original paper, a popular optimizer in a lot of the project, the original paper still chose to use SGD with momentum. Due to the fact that Adam fails to converge to an optimal solution in our specific setting, SGD with momentum does. Due to the following reasons that the original paper has already tuned the hyperparameters extensively for SGD optimization, we used the set  $\in [0.25, 0.1, 0.05, 0.01]$ , and used the momentum 0.9. We, on the other hand, used the same learning rates without extensively tuning them.

## Summary

This project is a reproduction work of one table in the paper "Striving for Simplicity: The All Convolutional Net". It compares four different kinds of branches. Results show that model A has a good consistency compared with model B, while model C differs. The influence of dropout and batch normalization to the results are also analyzed with several experiments. Using only batch normalization or combining batch normalization with dropout will not let the model converge.

## Discussion

From our results we obtained, All-CNN models performed much worse than the ConvPool and base CNN models. When we tried to use the validation, All-CNN models cannot converge to an optimal solution for model A and model B. The learning rate in the original paper will not let most models converge, so we decided to drop it and later used 0.001, we found out that smaller learning rates might not guarantee the fast converge, but it will generally converge.

## Further Reading

If you want to see more reproduction projects and even another researches on the same paper, please check the following links.

## Reference

Springenberg, Jost Tobias, et al. "Striving for simplicity: The all convolutional net." arXiv preprint arXiv:1511.02852 (2015).  
Li, Xiang, et al. "Understanding the disharmony between dropout and batch normalization by variance analysis." Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition. (2019).  
Ioffe, Sergey, and Christian Szegedy. "Batch normalization: Accelerating deep network training by reducing internal covariate shift." arXiv preprint arXiv:1502.03167 (2015).

## Appendix

- [Github Repository](#)
- [A Python Class to build all Models](#)
- [A Python Class for Training Procedure](#)
- [The Notebooks for Model A](#)
- [The Notebooks for Model B](#)

- [The Notebooks for Model C](#)
- [The Notebooks for Validation](#)
- [The Notebooks for DropOut and Batch Normalization](#)
- [The Notebooks for Optimizers](#)