

▼ Striving for Simplicity: The All Convolutional Net

Group 16: Wei-Tse Yang(5033292), Sunwei Wang(4345967), and Qingyuan Cao(5065232)

20/4/2020

In this notebook, we try to reproduce the TABLE 3 in the [original paper](#). The source code of one of the procedure on Google Colab can be found in [Github](#). **We adopt the original training procedure and change the models from scratch on Pytorch.**

Brief Introduction

The paper shows that replacing the max-pooling with the convolutional with increased strides can improve the performance. In this paper, we compare the models with max-pooling, models removing max-pooling, and models replacing max-pooling with convolutional layers with increased strides. The models replacing max-pooling with convolutional layers with increased strides generally have better performance. The authors tested 12 networks by designing 3 model bases and 3 branches.

Base: Model A, Model B, and Model C

Since the design of convolutional layers would influence the performance, the authors test three model bases. **Model A** uses the 5x5 strides but also adds one convolutional layer with 1x1 strides after that. **Model B** uses the 5x5 strides but also adds one convolutional layer with 1x1 strides after that. **Model C** uses the 3x3 strides.

Table 1: The three base networks used for classification on CIFAR-10 and CIFAR-100.

Model		
A	B	C
Input 32×32 RGB image		
5×5 conv. 96 ReLU	5×5 conv. 96 ReLU	3×3 conv. 96 ReLU
	1×1 conv. 96 ReLU	3×3 conv. 96 ReLU
3×3 max-pooling stride 2		
5×5 conv. 192 ReLU	5×5 conv. 192 ReLU	3×3 conv. 192 ReLU
	1×1 conv. 192 ReLU	3×3 conv. 192 ReLU
3×3 max-pooling stride 2		
3×3 conv. 192 ReLU		
1×1 conv. 192 ReLU		
1×1 conv. 10 ReLU		
global averaging over 6×6 spatial dimensions		
10 or 100-way softmax		

Branch: Model, Strided-CNN, ConvPool-CNN, and ALL-CNN

Each model base has one original model and three branches. “**Model**” is the model with max-pooling. “**Strided-CNN**” is the model replacing max-pooling with convolutional layers with increased strides. The better performance than “**Model**” and “**Strided-CNN**”. To solve it, “**ConvPool-CNN**” is proposed. “**ConvPool-CNN**” has one more convolutional layer before the pooling. “**ConvPool-CNN**” should have the same number of parameters as “**Model**”. “**All-CNN**” has a better performance than “**ConvPool-CNN**”, we can prove the better performance on “**All-CNN**” parameters. We show architecture with the base of model C in the following image.

Model		
Strided-CNN-C	ConvPool-CNN-C	All-CNN-C
Input 32×32 RGB image		
3×3 conv. 96 ReLU	3×3 conv. 96 ReLU	3×3 conv. 96 ReLU
3×3 conv. 96 ReLU	3×3 conv. 96 ReLU	3×3 conv. 96 ReLU
with stride $r = 2$	3×3 conv. 96 ReLU	
	3×3 max-pooling stride 2	3×3 conv. 96 ReLU
		with stride $r = 2$
3×3 conv. 192 ReLU	3×3 conv. 192 ReLU	3×3 conv. 192 ReLU
3×3 conv. 192 ReLU	3×3 conv. 192 ReLU	3×3 conv. 192 ReLU
with stride $r = 2$	3×3 conv. 192 ReLU	
	3×3 max-pooling stride 2	3×3 conv. 192 ReLU
		with stride $r = 2$

Experiment Setup

All 12 networks are trained on the CIFAR-10 with the stochastic gradient descent with a fixed momer rate γ is chosen from the set $\in [0.25, 0.1, 0.05, 0.01]$. It is also scheduled by multiplying with a fixed f . The paper only presents the best performance among all learning rates. Based on the source code, t CIFAR-10 test set. In other words, **the source code did not use the validation for hyper-parameter tu**

Reproduction Results

Our reproduction results are different from the paper as shown in the following table. First, we percei our reproduction, which is around 5~7%. Secondly, we obtain the same ranking in **Model A**, but we fa and **Model C**. This is because Model B and Model C are difficult to converge than Model A. Also, **ALL** the right place. During the training, the first three learning rates seem too large for the models. There rate=0.001 and a longer epoch=400 in **ALL-CNN-B**, where we add the * mark. Unfortunately, the mod

Model	Error Rate of Paper	Error Rate of Ours
Model A	12.47%	19.27%
Strided-CNN-A	13.46%	20.27%
ConvPool-CNN-A	10.21%	15.46%
ALL-CNN-A	10.30%	15.60%
Model B	10.20%	17.01%
Strided-CNN-B	10.98%	23.20%
ConvPool-CNN-B	9.33%	18.22%
ALL-CNN-B	9.10%	*29.48%
Model C	9.74%	13.07%
Strided-CNN-C	10.19%	15.49%
ConvPool-CNN-C	9.31%	14.39%
ALL-CNN-C	9.08%	17.89%

```
# The example of training procedure
# Choose the model A
training=Training(baseModel=[True,False,False])
```

```
# Create the dataset
training.createDataset()
# Choose the branch: All-CNN
training.modifiedModel=[False,False,False,True]
# Start Training
training.Procedure()
```

▼ Valudation

Since the source code did not use the validation set for hyper-parameters tuning, we conduct the validation on the training data to create the validation set. We show the results in the following table. The performance on the validation set is generally higher than the test error. However, if we compare the test error with the original test error, the test error is generally higher. The result reveals that the models might be overestimated.

Model	Validation Set Error	Test Error	Original Test Error
Model A	21.20%	20.45%	19.27%
Strided-CNN-A	21.72%	21.38%	20.27%
ConvPool-CNN-A	15.93%	17.04%	15.46%
ALL-CNN-A	17.57%	18.57%	15.60%
Model B	16.65%	17.81%	17.01%
Strided-CNN-B	17.53%	18.68%	23.20%
ConvPool-CNN-B	17.53%	17.51%	18.22%
ALL-CNN-B	*24.53%	*25.78%	*29.48%
Model C	14.13%	14.87%	13.07%
Strided-CNN-C	20.89%	21.67%	15.49%
ConvPool-CNN-C	17.81%	17.60%	14.39%
ALL-CNN-C	20.41%	19.13%	17.89%

```
# Validation can be conducted by settting validation equal to True
training=Training(validation=True,bestModel_allLR=True,baseModel=[True,False,False])
```

▼ DropOut and Batch Normalization

Dropout is a simple method to prevent neural networks from overfitting. In our all convolutional net p used to regularize all networks. Dropout was almost essential in all the state-of-the-art networks before batch normalization(BN). With the introduction of BN, it has shown its effectiveness and practicability in the evidence that when these two techniques are used combinedly in a network, the performance of the network is improved (Szegedy, 2015). In our study, we will investigate the results using BN and Dropout independently and Dropout simultaneously.

▼ BatchNorm *only*

```
model = Model(dropOut=False, BN=True)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 96, 30, 30]	7,296
BatchNorm2d-2	[-1, 96, 30, 30]	192
ReLU-3	[-1, 96, 30, 30]	0
MaxPool2d-4	[-1, 96, 14, 14]	0
Conv2d-5	[-1, 192, 12, 12]	460,992
BatchNorm2d-6	[-1, 192, 12, 12]	384
ReLU-7	[-1, 192, 12, 12]	0
MaxPool2d-8	[-1, 192, 5, 5]	0
Conv2d-9	[-1, 192, 5, 5]	331,968
BatchNorm2d-10	[-1, 192, 5, 5]	384
ReLU-11	[-1, 192, 5, 5]	0
Conv2d-12	[-1, 192, 5, 5]	37,056
BatchNorm2d-13	[-1, 192, 5, 5]	384
ReLU-14	[-1, 192, 5, 5]	0
Conv2d-15	[-1, 10, 5, 5]	1,930
BatchNorm2d-16	[-1, 10, 5, 5]	20
ReLU-17	[-1, 10, 5, 5]	0
AdaptiveAvgPool2d-18	[-1, 10, 1, 1]	0
Flatten-19	[-1, 10]	0
Total params: 840,606		

▼ BatchNorm with Dropout

```
model = Model(dropOut=True, BN=True)
```

Layer (type)	Output Shape	Param #
Dropout-1	[-1, 3, 32, 32]	0
Conv2d-2	[-1, 96, 30, 30]	7,296
BatchNorm2d-3	[-1, 96, 30, 30]	192
ReLU-4	[-1, 96, 30, 30]	0
MaxPool2d-5	[-1, 96, 14, 14]	0
Dropout-6	[-1, 96, 14, 14]	0
Conv2d-7	[-1, 192, 12, 12]	460,992
BatchNorm2d-8	[-1, 192, 12, 12]	384
ReLU-9	[-1, 192, 12, 12]	0
MaxPool2d-10	[-1, 192, 5, 5]	0
Dropout-11	[-1, 192, 5, 5]	0
Conv2d-12	[-1, 192, 5, 5]	331,968
BatchNorm2d-13	[-1, 192, 5, 5]	384
ReLU-14	[-1, 192, 5, 5]	0
Conv2d-15	[-1, 192, 5, 5]	37,056
BatchNorm2d-16	[-1, 192, 5, 5]	384
ReLU-17	[-1, 192, 5, 5]	0
Conv2d-18	[-1, 10, 5, 5]	1,930
BatchNorm2d-19	[-1, 10, 5, 5]	20
ReLU-20	[-1, 10, 5, 5]	0
AdaptiveAvgPool2d-21	[-1, 10, 1, 1]	0
Flatten-22	[-1, 10]	0
Total params: 840,606		

▼ Dropout *only*

```
model = Model(dropOut=True, BN=False)
```

Layer (type)	Output Shape	Param #
Dropout-1	[-1, 3, 32, 32]	0
Conv2d-2	[-1, 96, 30, 30]	7,296
ReLU-3	[-1, 96, 30, 30]	0
MaxPool2d-4	[-1, 96, 14, 14]	0
Dropout-5	[-1, 96, 14, 14]	0
Conv2d-6	[-1, 192, 12, 12]	460,992
ReLU-7	[-1, 192, 12, 12]	0
MaxPool2d-8	[-1, 192, 5, 5]	0
Dropout-9	[-1, 192, 5, 5]	0
Conv2d-10	[-1, 192, 5, 5]	331,968
ReLU-11	[-1, 192, 5, 5]	0
Conv2d-12	[-1, 192, 5, 5]	37,056
ReLU-13	[-1, 192, 5, 5]	0
Conv2d-14	[-1, 10, 5, 5]	1,930
ReLU-15	[-1, 10, 5, 5]	0
AdaptiveAvgPool2d-16	[-1, 10, 1, 1]	0
Flatten-17	[-1, 10]	0
Total params: 839,242		

Without using BatchNorm or Dropout

```
model = Model(dropOut=False, BN=False)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 96, 30, 30]	7,296
ReLU-2	[-1, 96, 30, 30]	0
MaxPool2d-3	[-1, 96, 14, 14]	0
Conv2d-4	[-1, 192, 12, 12]	460,992
ReLU-5	[-1, 192, 12, 12]	0
MaxPool2d-6	[-1, 192, 5, 5]	0
Conv2d-7	[-1, 192, 5, 5]	331,968
ReLU-8	[-1, 192, 5, 5]	0
Conv2d-9	[-1, 192, 5, 5]	37,056
ReLU-10	[-1, 192, 5, 5]	0
Conv2d-11	[-1, 10, 5, 5]	1,930
ReLU-12	[-1, 10, 5, 5]	0
AdaptiveAvgPool2d-13	[-1, 10, 1, 1]	0
Flatten-14	[-1, 10]	0
Total params: 839,242		

We compare the results of different combination of these two techniques and generated the table below

Model	BN only	BN + Dropout	Dropout only	No BN no Dropout
Model A	no converge	no converge	19.27%	13.82%

As shown in the table above, we used general Model A to study the two techniques Batch Normalization and Dropout. As we can see, adding Batch Normalization or Dropout increases or decreases our model performance in this case. We implemented BN layer between two

ReLU activations. Dropout is also applied in between convolution layers, and it is used after the pool used Dropout only, and we found out that using BN without Dropout or combining both BN with Dropout. Xiang, et al. (2019) stated in their papers that the worse performance may be the result of variance s lead the model to converge, giving the result of 19.27%. But the performance is still not as good as 1 might be due to that we did not have the time to tune hyperparameter, dropout rate, we only used the dropping out inputs and 50% otherwise.

▼ Optimizer

The default optimizer used in the paper and in our reproduction is Stochastic Gradient Descent (SGD) different optimizers since Adaptive Moment Estimation (Adam) is one of the most popular optimizat instead of SGD optimizer in our model. However, the model did not converge with Adam under the sp even though in theory Adam combines the advantage of two SGD variants, RMSProp and AdaGrad, it setting to converge to an optimal solution.

SGD with momentum optimizer:

```
self.optimizer = optim.SGD(model.parameters(), lr=lr, momentum=0.9)
```

Adam optimizer:

```
self.optimizer = optim.Adam(self.model.parameters(), lr=self.lr)
```

The optimizer is the main approach nowadays for training neural networks with minimizing its error i popular optimizer in a lot of the project, the original paper still chose to use SGD with momentum. D that Adam fails to converge to an optimal solution in our specific setting, SGD with momentum does due reasons that the original paper has already tuned the hyperparameters extensively for SGD optin the set $\in [0.25, 0.1, 0.05, 0.01]$, and used the momentum 0.9. We, on the other hand, used the same l without extensively tuning them.

Summary

This project is a reproduction work of one table in the paper "Striving for Simplicity: The All Convoluti four different kinds of branches. Results show that model A has a good consistency compared with C differ. The influence of dropout and batch normalization to the results are also analyzed with sever normalization only or combining batch normalization with dropout will not let the model converge.

Discussion

From our results we obtained, All-CNN models performed much worse than the ConvPool and base r we tried to use the validation, All-CNN models cannot converge to an optimal solution for model A ar learning rate in the original paper will not let most models converge, so we decided to drop it and late 0.001, we found out that smaller learning rates might not guarantee the fast converge, but it will gen

Further Reading

If you want to see more reproduction projects and even another researches on the same paper, please

Reference

Springenberg, Jost Tobias, et al. "Striving for simplicity: The all convolutional net." arXiv preprint arXiv:1511.02852 (2015).

Li, Xiang, et al. "Understanding the disharmony between dropout and batch normalization by variance on Computer Vision and Pattern Recognition. (2019).

Ioffe, Sergey, and Christian Szegedy. "Batch normalization: Accelerating deep network training by reducing internal covariate shift." arXiv preprint arXiv:1502.03167 (2015).

Appendix

- [Github Repository](#)
- [A Python Class to build all Models](#)
- [A Python Class for Training Procedure](#)
- [The Notebooks for Model A](#)
- [The Notebooks for Model B](#)
- [The Notebooks for Model C](#)
- [The Notebooks for Validation](#)
- [The Notebooks for DropOut and Batch Normalization](#)
- [The Notebooks for Optimizers](#)